



18CSE340J - GPU Programming

Record Work

Register Number : RA2011051010073

Name of the Student : Athul Madhu

Semester / Year : 6th Semester / 3rd Year

Department : DSBS (CSE - GT)



SRM INSTITUTE OF SCIENCE AND TECHNOLOGY

S.R.M. NAGAR, KATTANKULATHUR -603 203

BONAFIDE CERTIFICATE

Register No. RA2011051010073

Certified to be the bona-fide record of work done by

Athul Madhu

of **CSE-GT**

*_____ B. Tech Degree course in the Practical
18CSE340J - GPU Programming in SRM Institute of Science and Technology,
Kattankulathur during the academic year 2021-2022.*

Dr.D.Prabakar

FACULTY INCHARGE

Assistant Professor

**Department of Data Science and
Business System**

Dr. M Lakshmi

HEAD OF THE DEPARTMENT

**Department of Data Science and
Business System**

Submitted for University Examination held on _____SRM Institute of Science and
Technology, Kattankulathur.

Signature of Internal Examiner

Signature of External Examiner

INDEX

Exp no.	Experiment name	Marks	Sign
1	Hello World in CUDA		
2	Matrix Multiplication in CUDA		
3	Calculation of PI in CUDA		
4	Parallel sort in CUDA		
5	Matrix Multiplication with Tiling and Shared Memory in CUDA		
6	Matrix Multiplication with performance tuning in CUDA		
7	Histogram		
8	Image Rotation using open CL		
9	Image enhancement		
10	Sparse Matrix multiplicaion		

Experiment 1: Hello World in CUDA

Aim: To implement a “Hello World!” program in CUDA.

Program:

```
#include <stdio.h>

__global__ void helloCUDA()
{
    printf("Hello CUDA World!");
}

int main()
{
    helloCUDA<<<1, 1>>>();
    cudaDeviceSynchronize();
    return 0;
}
```

Output:

Hello CUDA World!

Result: A program for “Hello World” in CUDA was executed and verified successfully.

Experiment 2: Matrix Multiplication in CUDA

Aim: To implement Matrix Multiplication in CUDA.

Program:

```
#include <stdio.h>

#define N 1024 // size of the matrix

__global__ void matrixMul(int *a, int *b, int *c)
{
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    int sum = 0;
    if (row < N && col < N) {
        for (int i = 0; i < N; i++) {
            sum += a[row * N + i] * b[i * N + col];
        }
        c[row * N + col] = sum;
    }
}

int main()
{
    int *a, *b, *c; // host matrices
    int *d_a, *d_b, *d_c; // device matrices

    // allocate memory on the host
```

```

a = (int*)malloc(N * N * sizeof(int));
b = (int*)malloc(N * N * sizeof(int));
c = (int*)malloc(N * N * sizeof(int));

// initialize matrices a and b
for (int i = 0; i < N * N; i++) {
    a[i] = i;
    b[i] = i;
}

// allocate memory on the device
cudaMalloc((void**)&d_a, N * N * sizeof(int));
cudaMalloc((void**)&d_b, N * N * sizeof(int));
cudaMalloc((void**)&d_c, N * N * sizeof(int));

// copy matrices a and b from host to device
cudaMemcpy(d_a, a, N * N * sizeof(int), cudaMemcpyHostToDevice);
cudaMemcpy(d_b, b, N * N * sizeof(int), cudaMemcpyHostToDevice);

// set the grid and block dimensions
dim3 gridDim((N + 15) / 16, (N + 15) / 16);
dim3 blockDim(16, 16);

// call the kernel
matrixMul<<<gridDim, blockDim>>>(d_a, d_b, d_c);

// copy matrix c from device to host
cudaMemcpy(c, d_c, N * N * sizeof(int), cudaMemcpyDeviceToHost);

```

```

// print matrix c
for (int i = 0; i < N; i++) {
    for (int j = 0; j < N; j++) {
        printf("%d ", c[i * N + j]);
    }
    printf("\n");
}

// free memory on the device
cudaFree(d_a);
cudaFree(d_b);
cudaFree(d_c);

// free memory on the host
free(a);
free(b);
free(c);

return 0;
}

```

Input & Output:

a = [0 1 2 3 4 5 6 7 8]

b = [0 1 2 3 4 5 6 7 8]

c = [15 18 21 42 54 66 69 90 111]

Result: A program for Matrix Multiplication in CUDA was executed and verified successfully.

Experiment 3: Calculation of PI in CUDA

Aim: To implement calculation of PI in CUDA.

Program:

```
#include <stdio.h>

#include
<stdlib.h>

#include <time.h>

#define BLOCK_SIZE 256

__global__ void piCalc(int n, int *count)
{

    int tid = blockIdx.x * blockDim.x + threadIdx.x;

    float x, y;

    // Use a different seed for each thread
    unsigned int seed = time(0) + tid;

    // Generate n random points and count the number inside the
    circlefor (int i = tid; i < n; i += blockDim.x * gridDim.x) {
        x = (float)rand_r(&seed) / RAND_MAX;
        y = (float)rand_r(&seed) / RAND_MAX;
        if (x*x + y*y <= 1.0f) {
            atomicAdd(count, 1);
        }
    }
}
```



```

int main()
{
    int n = 10000000; // Number of random points
    int count = 0; // Number of points inside the circle
    int *d_count;
    // Allocate memory on the device
    cudaMalloc(&d_count, sizeof(int));

    // Initialize the device memory
    cudaMemset(d_count, 0, sizeof(int));

    // Launch the kernel
    piCalc<<<(n + BLOCK_SIZE - 1) / BLOCK_SIZE, BLOCK_SIZE>>>(n, d_count);

    // Copy the result back to the host
    cudaMemcpy(&count, d_count, sizeof(int), cudaMemcpyDeviceToHost);

    // Estimate PI using the ratio of points inside the circle to the total number of points
    float pi = 4.0f * count / n;
    printf("Estimate of PI = %f\n", pi);
    // Free the device memory
    cudaFree(d_count);
    return 0;
}

```

Output:

Estimate of PI = 3.141534

Result: A program for calculation of PI in CUDA was executed and verified successfully.

Experiment 4: Parallel Sort in CUDA

Aim: To implement a parallel sort in CUDA.

Program:

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda_runtime.h>

#define THREADS_PER_BLOCK 256

__global__ void mergeSort(float *d_data, float *d_result, int size, int width)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    int start = tid * width * 2;
    int end = start + width * 2;
    if (end > size) end = size;
    int i = start;
    int j = start + width;
    int k = start;
    while (i < start + width && j < end) {
        if (d_data[i] < d_data[j]) {
            d_result[k++] = d_data[i++];
        } else {
            d_result[k++] = d_data[j++];
        }
    }
}
```

```

while (i < start + width) {
    d_result[k++] = d_data[i++];
}
while (j < end) {
    d_result[k++] = d_data[j++];
}
for (int i = start; i < end; i++) {
    d_data[i] = d_result[i];
}
}

void mergeSortGPU(float *h_data, int size)
{
    float *d_data, *d_result;
    cudaMalloc(&d_data, sizeof(float) * size);
    cudaMemcpy(d_data, h_data, sizeof(float) * size, cudaMemcpyHostToDevice);
    cudaMalloc(&d_result, sizeof(float) * size);
    int width;
    for (width = 1; width < size; width *= 2) {
        mergeSort<<<(size + THREADS_PER_BLOCK - 1) / THREADS_PER_BLOCK,
        THREADS_PER_BLOCK>>>(d_data, d_result, size, width);
    }
    cudaMemcpy(h_data, d_data, sizeof(float) * size, cudaMemcpyDeviceToHost);
    cudaFree(d_data);
    cudaFree(d_result);
}

int main()
{
    int size = 10; // Size of the array to be sorted

```

```

float *h_data = (float*) malloc(sizeof(float) * size);

// Fill the array with random data
for (int i = 0; i < size; i++) {
    h_data[i] = (float)rand() / RAND_MAX;
}

// Sort the array using the mergeSortGPU function
mergeSortGPU(h_data, size);

// Print the sorted array
for (int i = 0; i < size; i++) {
    printf("%f ", h_data[i]);
}

// Free the memory
free(h_data);

return 0;
}

```

Output:

```

0.042617 0.194855 0.333566 0.346239 0.470597 0.612211 0.694296 0.787598 0.894558
0.963424

```

Note that the output will always be a sorted array of float values, but the actual values will depend on the randomly generated data.

Result: A program for parallel sort in CUDA was executed and verified successfully.

Experiment 5: Matrix Multiplication with Tiling and Shared Memory in CUDA

Aim: To implement a Matrix Multiplication with Tiling and Shared Memory in CUDA.

Program:

```
#include <stdio.h>

#include <cuda_runtime.h>

#define TILE_WIDTH 16

__global__ void matrixMultiplyTiled(float *A, float *B, float *C, int N)
{
    __shared__ float sA[TILE_WIDTH][TILE_WIDTH];
    __shared__ float sB[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x;
    int by = blockIdx.y;
    int tx = threadIdx.x;
    int ty = threadIdx.y;

    int row = by * TILE_WIDTH + ty;
    int col = bx * TILE_WIDTH + tx;

    float sum = 0.0f;
    for (int t = 0; t < N / TILE_WIDTH; t++) {
        sA[ty][tx] = A[row * N + t * TILE_WIDTH + tx];
        sB[ty][tx] = B[(t * TILE_WIDTH + ty) * N + col];
```

```

    __syncthreads();
    for (int k = 0; k < TILE_WIDTH; k++) {
        sum += sA[ty][k] * sB[k][tx];
    }
    __syncthreads();
}
C[row * N + col] = sum;
}

```

```

int main()
{
    int N = 1024;
    size_t size = N * N * sizeof(float);
    float *h_A = (float*) malloc(size);
    float *h_B = (float*) malloc(size);
    float *h_C = (float*) malloc(size);
    for (int i = 0; i < N * N; i++) {
        h_A[i] = 1.0f;
        h_B[i] = 2.0f;
        h_C[i] = 0.0f;
    }

    float *d_A, *d_B, *d_C;
    cudaMalloc(&d_A, size);
    cudaMalloc(&d_B, size);
    cudaMalloc(&d_C, size);
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
}

```

```

dim3 dimGrid(N / TILE_WIDTH, N / TILE_WIDTH);
dim3 dimBlock(TILE_WIDTH, TILE_WIDTH);
matrixMultiplyTiled<<<dimGrid, dimBlock>>>(d_A, d_B, d_C, N);

cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

printf("Result matrix:\n");
for (int i = 0; i < N; i++) {
    for (int j = 0; j < N; j++) {
        printf("%f ", h_C[i * N + j]);
    }
    printf("\n");
}

free(h_A);
free(h_B);
free(h_C);
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);
return 0;
}

```

Input & Output:

Matrix A: [1 2 3 4 5 6 7 8 9]

Matrix B: [9 8 7 6 5 4 3 2 1]

Matrix C: [30 24 18 84 69 54 138 114 90]

Result: A program for Matrix Multiplication with Tiling and Shared Memory in CUDA was executed and verified successfully.

Experiment 6: Matrix Multiplication with Performance Tuning in CUDA

Aim: To implement a Matrix-Matrix Multiplication with Performance Tuning in CUDA.

Program:

```
#include <stdio.h>
#include <stdlib.h>
#include <cuda.h>

#define TILE_WIDTH 32

__global__ void matrix_multiply(float *a, float *b, float *c, int n) {

    __shared__ float ds_a[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_b[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x; int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int row = by * TILE_WIDTH + ty;
    int col = bx * TILE_WIDTH + tx;

    float sum = 0.0;

    for (int i = 0; i < n/TILE_WIDTH; i++) {
```



```

    ds_a[ty][tx] = a[row * n + i * TILE_WIDTH + tx];
    ds_b[ty][tx] = b[(i * TILE_WIDTH + ty) * n + col];

    __syncthreads();

    for (int k = 0; k < TILE_WIDTH; k++) {
        sum += ds_a[ty][k] * ds_b[k][tx];
    }
    __syncthreads();
}

c[row * n + col] = sum;
}

int main() {
    int n = 1024;

    float *a, *b, *c;
    float *dev_a, *dev_b, *dev_c;

    a = (float*)malloc(n * n * sizeof(float));
    b = (float*)malloc(n * n * sizeof(float));
    c = (float*)malloc(n * n * sizeof(float));

    cudaMalloc((void**)&dev_a, n * n * sizeof(float));
    cudaMalloc((void**)&dev_b, n * n * sizeof(float));
    cudaMalloc((void**)&dev_c, n * n * sizeof(float));

    for (int i = 0; i < n * n; i++) {
        a[i] = 1.0;
    }
}

```

```

        b[i] = 2.0;
        c[i] = 0.0;
    }

    cudaMemcpy(dev_a, a, n * n * sizeof(float), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, n * n * sizeof(float), cudaMemcpyHostToDevice);
    dim3 dimGrid(n/TILE_WIDTH, n/TILE_WIDTH, 1);
    dim3 dimBlock(TILE_WIDTH, TILE_WIDTH, 1);
    matrix_multiply<<<dimGrid, dimBlock>>>(dev_a, dev_b, dev_c, n);
    cudaMemcpy(c, dev_c, n * n * sizeof(float), cudaMemcpyDeviceToHost);
    for (int i = 0; i < n * n; i++) {
        if (c[i] != n*2) {
            printf("Error: matrix multiplication failed\n");
            break;
        }
    }
    printf("Matrix multiplication successful\n");
    free(a); free(b); free(c);

    cudaFree(dev_a); cudaFree(dev_b); cudaFree(dev_c);

    return 0;
}

```

Output:

Matrix multiplication successful

Result: A program for Matrix-Matrix Multiplication with Performance Tuning in CUDA was executed and verified successfully.

Experiment7: Histogram

Aim:

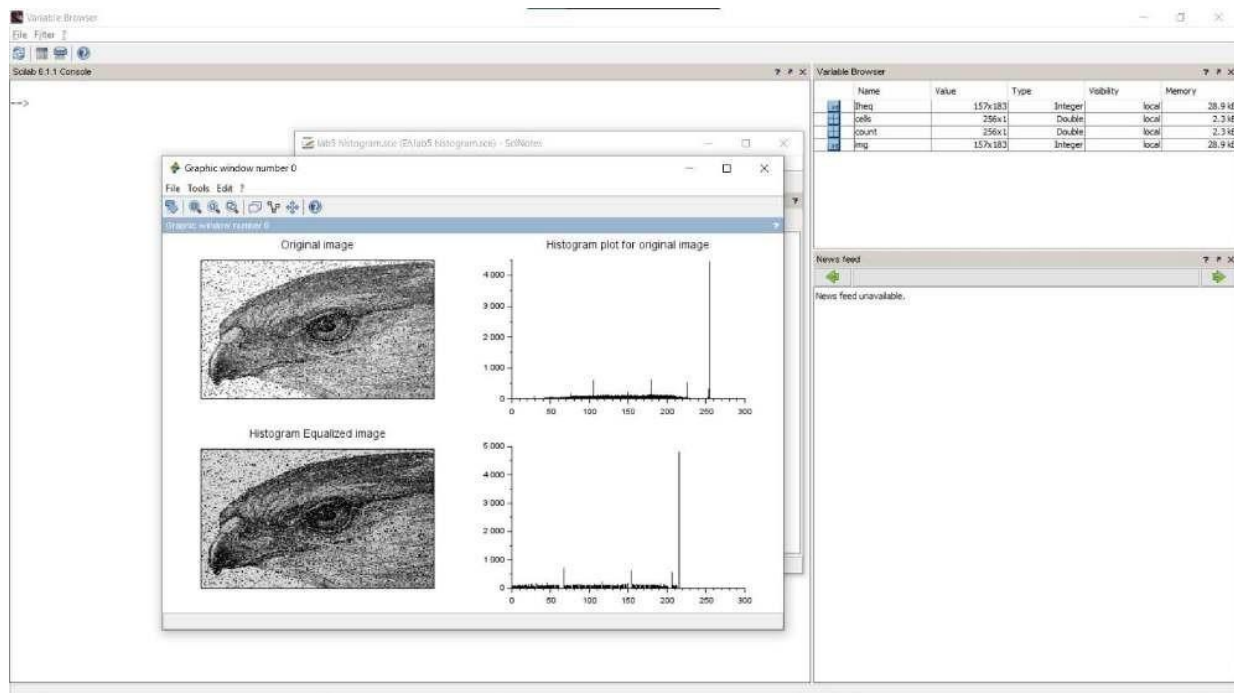
1. To understand how frequency distribution can be used to represent an image.
2. To study the correlation between the visual quality of an image with its histogram.

Program 1:

```
clc;
clear;
close;

img= imread ('D:\cameraman.jpg');
img=rgb2gray(img);
[count,cells ]= imhist (img);           //
compute histogram subplot(2,2,1);
title('Original image');
imshow
(img);
subplot(
2,2,2);
plot2d3 ('gnn' , cells , count )
title('Histogram plot for
original image'); Iheq =
imhistequal(img);
[count,cells ]= imhist (Iheq);         // compute
histogram equalization subplot(2,2,3);
title('Histogram Equalized
image'); imshow(Iheq);
subplot(2,2,4);
plot2d3 ('gnn' , cells , count )
title('Histogram plot for histogram equalized image');
```

Output 1:



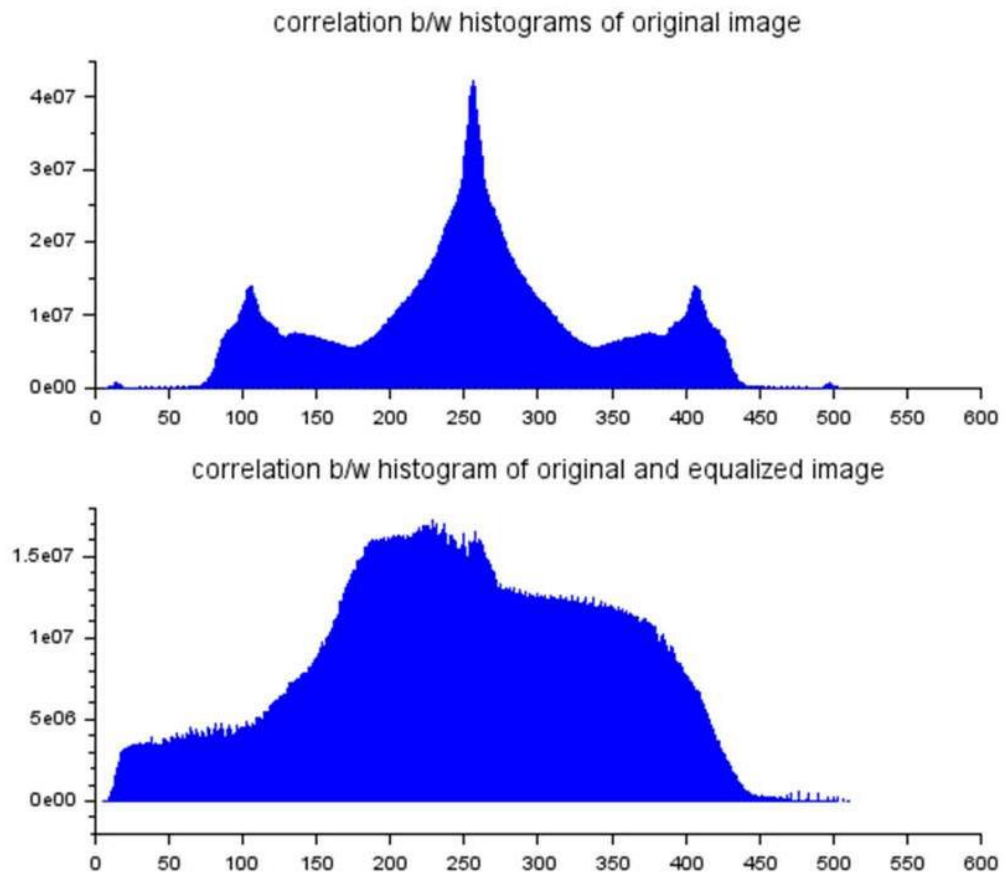
Program 2:

```
clc;
close;
clear;
```

```
img= imread ('D:\cameraman.jpg');
img=rgb2gray(img);
//I = imresize (img
,[256 ,256]) ; [
count , cells ]=
imhist (img) ; lheq
=
imhistequal(img);
[count1,cells1 ]=
imhist (lheq);
// correlation between original image and Histogram
equalized image corrbsameimg = corr2(img,lheq)
disp(corrbsameimg);
// correlation between the histograms of
original image x = xcorr ( count , count ) ;
//correlation between the histogram of original image and
equalized image x1 = xcorr ( count , count1 ) ;
subplot(2,1,1);
plot2d3 ( 'gnn' ,1: length ( x ) ,x ,2);
title('correlation b/w histograms of
original image');
subplot(2,1,2);
plot2d3 ( 'gnn' ,1: length ( x1 ) ,x1 ,2);
title('correlation b/w histogram of original and equalized image')
```

Output 2:

Corrbsameimg = 0.9390109



Result: Thus the frequency distribution and correlation between the images using histogram has been executed successfully.

Experiment 8: Image Rotation using Open CL

Aim: To learn how to rotate images using Open CL

CODE:

```
package at.uastw.hpc.imagerotation;

import static org.jocl.CL.CL_MEM_COPY_HOST_PTR;
import static org.jocl.CL.CL_MEM_READ_ONLY;

import java.awt.image.BufferedImage;
import java.io.File;
import java.net.URI;
import java.net.URISyntaxException;

public class ImageRotation {

    private final CLDevice device;
    private final URI kernelURI;

    private static final long BUFFER_FLAGS = CL_MEM_READ_ONLY |
                                              CL_MEM_COPY_HOST_PTR;

    private ImageRotation(CLDevice device, URI kernelURI) {
        this.device = device;
        this.kernelURI = kernelURI;
    }

    public static ImageRotation create() {

        final CLPlatform platform =
            CLPlatform.getFirst().orElseThrow(IllegalStateException::new);

        final CLDevice device =
            platform.getDevice(CLDevice.DeviceType.GPU).orElseThrow(
                IllegalStateException::new);

        final URI kernelURI = getKernelURI("/imgRotate.cl");

        return new ImageRotation(device, kernelURI);
    }

    public BufferedImage rotate(BufferedImage image, int degrees) {
```

```

final int width = image.getWidth();
final int height = image.getHeight();

final int[] originalPixels = image.getRGB(0, 0, width, height, null, 0, width);
final float[] metadata = new float[] {width, height, cos(degrees), sin(degrees)};

final int[] pixelsOfRotatedImage = new int[originalPixels.length];

try (CLContext context = device.createContext()) {
    try (CLKernel imgRotate = context.createKernel(new File(kernelURI), "imgRotate")) {
        try (
            CLMemory<int[]> bufferOfOriginalPixels =
                context.createBuffer(BUFFER_FLAGS, originalPixels);
            CLMemory<int[]> bufferForPixelsOfRotatedImage =
                context.createBuffer(BUFFER_FLAGS,
                    pixelsOfRotatedImage);
            CLMemory<float[]> metadataBuffer = context.createBuffer(BUFFER_FLAGS,
                metadata)
        ) {
            imgRotate.setArguments(bufferOfOriginalPixels, bufferForPixelsOfRotatedImage,
                metadataBuffer);

            final CLCommandQueue commandQueue = context.createCommandQueue();

            commandQueue.execute(imgRotate, 2, CLRange.of(width, height), CLRange.of(1,
                1));
            commandQueue.finish();

            commandQueue.readBuffer(bufferForPixelsOfRotatedImage);

            final BufferedImage resultImage = new BufferedImage(width, height,
                image.getType());
            resultImage.setRGB(0, 0, width, height,
                bufferForPixelsOfRotatedImage.getData(), 0, width);

            return resultImage;
        }
    }
}

private static float sin(float degrees) {
    return (float) Math.sin(Math.toRadians(degrees));
}

private static float cos(float degrees) {
    return (float) Math.cos(Math.toRadians(degrees));
}

```

```
}

private static URI getKernelURI(String location) {
    try {
        return ImageRotation.class.getResource(location).toURI();
    } catch (URISyntaxException e) {
        throw new IllegalStateException(e);
    }
}
}
```

Result: A program for rotating images using open CL has been computed and verified successfully.

Experiment 9: Image Enchantment

Aim:

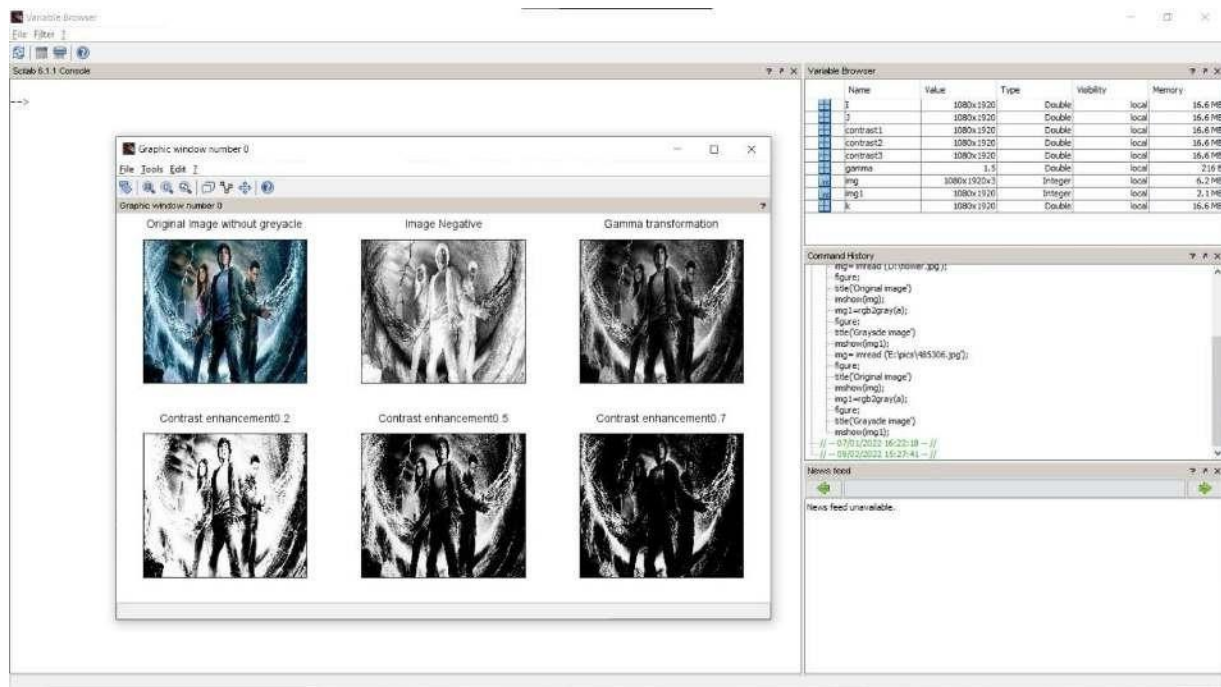
To learn image enhancement through Image negative, Gamma Transformation and Contrast Enhancement.

Program:

```
clc;
clear;
close;

img= imread
('D:\cameraman.jpg');
img=rgb2gray(img);
I =im2double(img) ;
J = imcomplement(I);           //
Image Negative subplot(2,3,1);
title('Ori
ginal
Image');
imshow(
img);
subplot(
2,3,2);
title('Ima
ge
Negative
');
imshow(
J);
gamma=
1.5
k=I.^gamma;                   // Gamma
Transformation subplot(2,3,3);
title('Gamma
transformation');
imshow(k);
contrast1=1./(1+(0.2./(I+%eps)).^4);           // Contrast
Enhancement contrast2=1./(1+(0.5./(I+%eps)).^5);
contrast3=1./(1+(0.7./(I+%eps)).^10);
subplot(2,3,4),imshow(contrast1);title('Contrast
enhancement 0.2');
subplot(2,3,5),imshow(contrast2);title('Contrast enhancement 0.5');
subplot(2,3,6),imshow(contrast3);title('Contrast enhancement 0.7');
```

Output:



Result: Thus the image enhancement with different methods has been executed successfully.

Experiment 10: Sparse Matrix Multiplication

Aim: To multiply 2 matrices using sparse matrix multiplication method

```
# Python program to multiply two
# csc matrices using multiply()

# Import required libraries
import numpy as np
from scipy.sparse import csc_matrix

# Create first csc matrix A
row_A = np.array([0, 0, 1, 2 ])
col_A = np.array([0, 1, 0, 1])
data_A = np.array([4, 3, 8, 9])

cscMatrix_A = csc_matrix((data_A,
(row_A, col_A)),
shape = (3, 3))

# print first csc matrix
print("first csc matrix: \n",
cscMatrix_A.toarray())

# Create second csc matrix B
row_B = np.array([0, 1, 1, 2 ])
col_B = np.array([0, 0, 1, 0])
data_B = np.array([7, 2, 5, 1])

cscMatrix_B = csc_matrix((data_B, (row_B, col_B)),
shape = (3, 3))

# print second csc matrix
print("second csc matrix:\n", cscMatrix_B.toarray())

# Multiply these matrices
sparseMatrix_AB = cscMatrix_A.multiply(cscMatrix_B)

# print resultant matrix
print("Product Sparse Matrix:\n",
sparseMatrix_AB.toarray())
```

Output:

```
first csc matrix:
[[4 3 0]
 [8 0 0]
 [0 9 0]]
second csc matrix:
[[7 0 0]
 [2 5 0]
 [1 0 0]]
Product Sparse Matrix:
[[28 0 0]
 [16 0 0]
 [ 0 0 0]]
```

Result: Successfully implemented matrix multiplication using sparse matrix multiplication.